


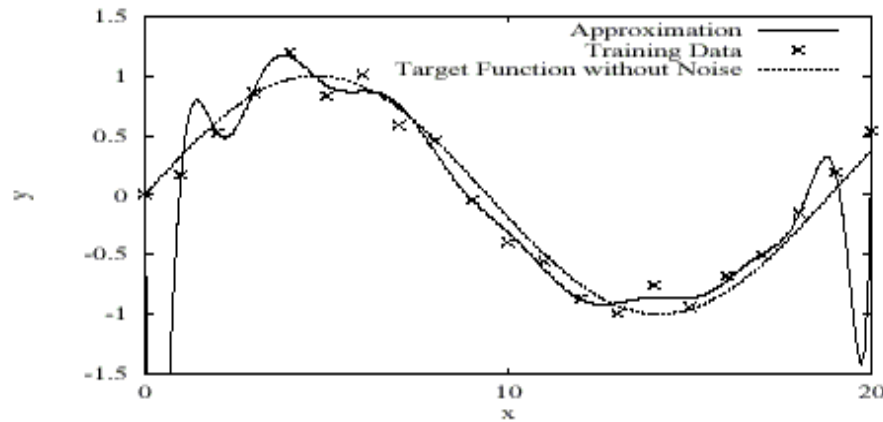
Inductive Bias: How to generalize on novel data

Overfitting


Noise vs. Exceptions




Order 2



Order 16



Order 10




Order 20

Non-Linear Tasks

- Linear Regression will not generalize well to the task below
- Needs a non-linear surface
- Could do a feature pre-process as with the quadric machine
 - For example, we could use an arbitrary polynomial in x
 - Thus it is still linear in the coefficients, and can be solved with delta rule, etc.

$$Y = \beta_0 + \beta_1 X + \beta_2 X^2 + \dots + \beta_n X^n$$

- What order polynomial should we use? – Overfit issues can occur



Regression Regularization

- How to avoid overfit – Keep the model simple
 - For regression, keep the function smooth
 - Inductive bias is that $f(x) \approx f(x \pm \varepsilon)$
- Regularization approach: $F(h) = \text{Error}(h) + \lambda \cdot \text{Complexity}(h)$
 - Tradeoff accuracy vs complexity
- Ridge Regression – Minimize:
 - $F(\mathbf{w}) = \text{TSS}(\mathbf{w}) + \lambda \|\mathbf{w}\|^2 = \sum (\text{predicted}_i - \text{actual}_i)^2 + \lambda \sum w_i^2$
 - Gradient of $F(\mathbf{w})$: $\Delta w_i = c(t - \text{net})x_i - \lambda w_i$ (Weight decay)
 - Especially useful when the features are a non-linear transform from the initial features (e.g. polynomials in x)
 - Also when the number of initial features is greater than the number of examples
 - Lasso regression uses an L1 vs an L2 weight penalty: $\text{TSS}(\mathbf{w}) + \lambda \sum |w_i|$

Hypothesis Space

- The Hypothesis space H is the set all the possible models h which can be learned by the current learning algorithm
 - e.g. Set of possible weight settings for a perceptron
- Restricted hypothesis space
 - Can be easier to search
 - May avoid overfit since they are usually simpler (e.g. linear or low order decision surface)
 - Often will underfit
- Unrestricted Hypothesis Space
 - Can represent any possible function and thus can fit the training set well
 - Mechanisms must be used to avoid overfit


Avoiding Overfit - Regularization

- Regularization: *any modification we make to a learning algorithm that is intended to reduce its generalization error but not its training error*
- Occam's Razor – William of Ockham (c. 1287-1347)
- Simplest accurate model: accuracy vs. complexity trade-off. Find $h \in H$ which minimizes an objective function of the form:

$$F(h) = \text{Error}(h) + \lambda \cdot \text{Complexity}(h)$$

- Complexity could be number of nodes, size of tree, magnitude of weights, order of decision surface, etc. L2 and L1 common.
- More Training Data (vs. overtraining on same data)
 - Also Data set augmentation – Fake data, Can be very effective, Jitter, but take care...
 - Denoising – add random noise to inputs during training – can act as a regularizer
 - Adding noise to nodes, weights, outputs, etc. E.g. Dropout (discuss with ensembles)
- Most common regularization approach: *Early Stopping* – Start with simple model (small parameters/weights) and stop training as soon as we attain good generalization accuracy (before parameters get large)
 - Use a validation Set (next slide: requires separate test set)
- Will discuss other approaches with specific models

Stopping/Model Selection with Validation Set



- There is a different model h after each epoch
- Select a model in the area where the validation set accuracy flattens
 - When no improvement occurs over m epochs
- The validation set comes out of training set data
- Still need a separate test set to use after selecting model h to predict future accuracy
- Simple and unobtrusive, does not change objective function, etc
 - Can be done in parallel on a separate processor
 - Can be used alone or in conjunction with other regularizers

Inductive Bias

- The approach used to decide how to generalize novel cases
- One common approach is Occam's Razor – The *simplest* hypothesis which *explains/fits* the data is usually the best
- Many other rationale biases and variations

$$ABC \Rightarrow Z$$

$$A\bar{B}C \Rightarrow Z$$

$$AB\bar{C} \Rightarrow Z$$

$$A\bar{B}\bar{C} \Rightarrow Z$$

$$\bar{A}\bar{B}\bar{C} \Rightarrow \bar{Z}$$

$$\bar{A}BC \Rightarrow ?$$

- When you get the new input $\bar{A} B C$. What is your output?

One Definition for Inductive Bias

Inductive Bias: Any basis for choosing one generalization over another, other than strict consistency with the observed training instances

Sometimes just called the *Bias* of the algorithm (don't confuse with the bias weight in a neural network).

Bias-Variance Trade-off – Will discuss in more detail when we discuss ensembles

Some Inductive Bias Approaches

- Restricted Hypothesis Space - Can just try to minimize error since hypotheses are already simple
 - Linear or low order threshold function
 - k -DNF, k -CNF, etc.
 - Low order polynomial
- Preference Bias – Prefer one hypothesis over another even though they have similar training accuracy
 - Occam's Razor
 - “Smallest” DNF representation which matches well
 - Shallow decision tree with high information gain
 - Neural Network with low validation error and small magnitude weights

Need for Bias

2^{2^n} Boolean functions of n inputs

<u>x1</u>	<u>x2</u>	<u>x3</u>	<u>Class</u>	<u>Possible Consistent Function Hypotheses</u>
0	0	0	1	
0	0	1	1	
0	1	0	1	
0	1	1	1	
1	0	0		
1	0	1		
1	1	0		
1	1	1	?	

Need for Bias

2^{2^n} Boolean functions of n inputs

<u>x1</u>	<u>x2</u>	<u>x3</u>	<u>Class</u>	<u>Possible Consistent Function Hypotheses</u>
0	0	0	1	1
0	0	1	1	1
0	1	0	1	1
0	1	1	1	1
1	0	0		0
1	0	1		0
1	1	0		0
1	1	1	?	0

Need for Bias

2^{2^n} Boolean functions of n inputs

<u>x1</u>	<u>x2</u>	<u>x3</u>	<u>Class</u>	<u>Possible Consistent Function Hypotheses</u>
0	0	0	1	1 1
0	0	1	1	1 1
0	1	0	1	1 1
0	1	1	1	1 1
1	0	0		0 0
1	0	1		0 0
1	1	0		0 0
1	1	1	?	0 1

Need for Bias

2^{2^n} Boolean functions of n inputs

x1	x2	x3	Class	<u>Possible Consistent Function Hypotheses</u>													
0	0	0	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1
0	0	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1
0	1	0	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1
0	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1
1	0	0		0	0	0	0	0	0	0	1	1	1	1	1	1	1
1	0	1		0	0	0	0	1	1	1	1	0	0	0	0	1	1
1	1	0		0	0	1	1	0	0	1	1	0	0	1	1	0	1
1	1	1	?	0	1	0	1	0	1	0	1	0	1	0	1	0	1

Without an Inductive Bias we have no rationale to choose one hypothesis over another and thus a random guess would be as good as any other option.

Need for Bias

2^{2^n} Boolean functions of n inputs

x1	x2	x3	Class	Possible Consistent Function Hypotheses													
0	0	0	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1
0	0	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1
0	1	0	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1
0	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1	1
1	0	0		0	0	0	0	0	0	0	1	1	1	1	1	1	1
1	0	1		0	0	0	0	1	1	1	1	0	0	0	0	1	1
1	1	0		0	0	1	1	0	0	1	1	0	0	1	1	0	1
1	1	1	?	0	1	0	1	0	1	0	1	0	1	0	1	0	1

Inductive Bias guides which hypothesis we should prefer?

What happens in this case if we use simplicity (Occam's Razor) as our inductive Bias?

Learnable Problems


- “Raster Screen” Problem
- Pattern Theory
 - Regularity in a task
 - Compressibility
- Don’t care features and Impossible states
- Interesting/Learnable Problems
 - What we actually deal with
 - Can we formally characterize them?
- Learning a training set vs. generalizing
 - A function where each output is set randomly (coin-flip)
 - Output class is independent of all other instances in the data set
- Computability vs. Learnability (Optional)

Computability and Learnability – Finite Problems

- Finite problems assume finite number of mappings (Finite Table)
 - Fixed input size arithmetic
 - Random memory in a RAM
- Learnable: Can do better than random on novel examples

Computability and Learnability – Finite Problems

- Finite problems assume finite number of mappings (Finite Table)
 - Fixed input size arithmetic
 - Random memory in a RAM
- Learnable: Can do better than random on novel examples

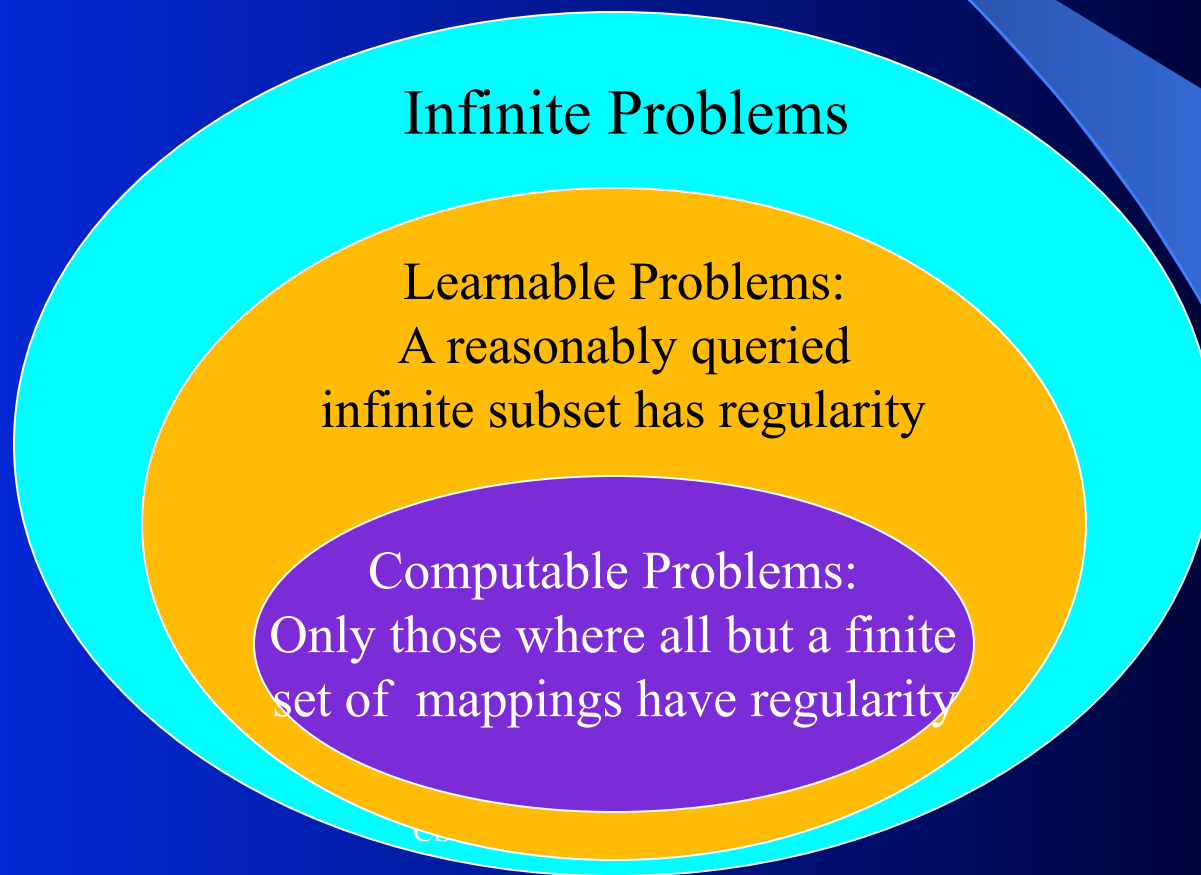


Computability and Learnability – Infinite Problems

- Infinite number of mappings (Infinite Table)
 - Arbitrary input size arithmetic
 - Halting Problem (no limit on input size)
 - Do two arbitrary strings match

Computability and Learnability – Infinite Problems

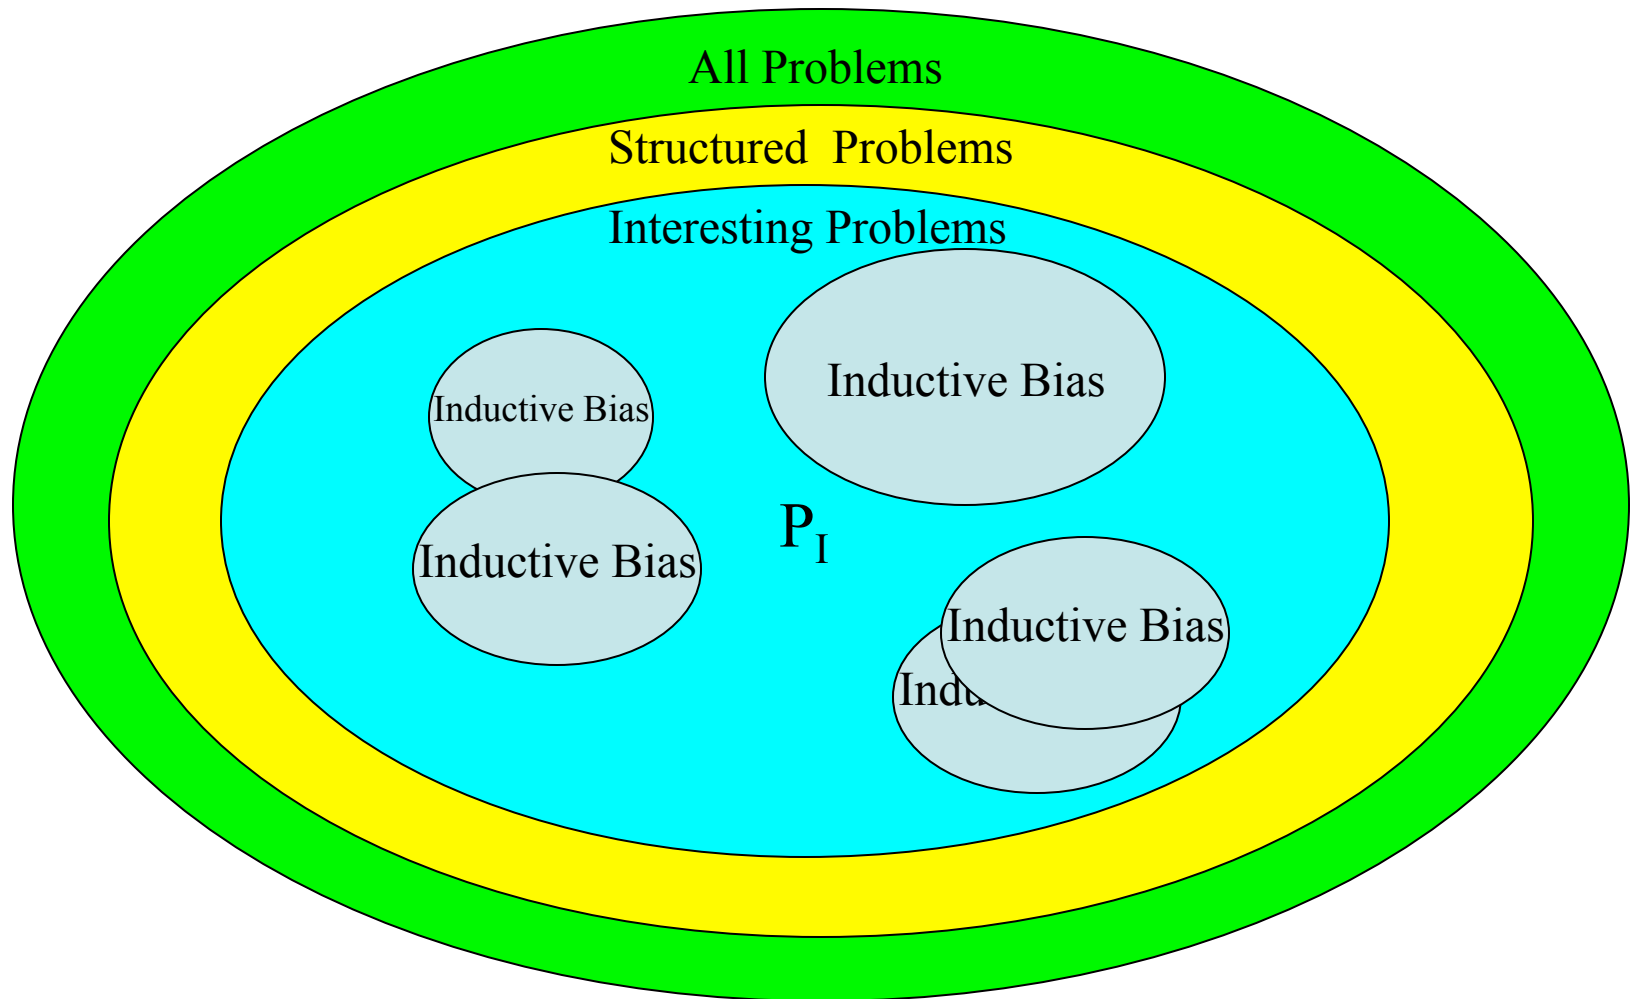
- Infinite number of mappings (Infinite Table)
 - Arbitrary input size arithmetic
 - Halting Problem (no limit on input size)
 - Do two arbitrary strings match



No Free Lunch

- Any inductive bias chosen will have equal accuracy compared to any other bias over *all* possible functions/tasks, assuming all functions are equally likely. If a bias is correct on some cases, it must be incorrect on equally many cases.
- Is this a problem?
 - Random vs. Regular
 - Anti-Bias? (even though regular)
 - The “Interesting” Problems – subset of learnable?
- Are all functions equally likely in the real world?

Interesting Problems and Biases



More on Inductive Bias

- Inductive Bias requires some set of prior assumptions about the tasks being considered and the learning approaches available
- Tom Mitchell's definition: Inductive Bias of a learner is the set of additional assumptions sufficient to justify its inductive inferences as deductive inferences
- We consider standard ML algorithms/hypothesis spaces to be different inductive biases: C4.5 (Greedy best attributes), Backpropagation (simple to complex), etc.

Which Bias is Best?

- Not one Bias that is best on all problems
- Our experiments
 - Over 50 real world problems
 - Over 400 inductive biases – mostly variations on critical variable biases vs. similarity biases
- Different biases were a better fit for different problems
- Given a data set, which Learning model (Inductive Bias) should be chosen?

Automatic Discovery of Inductive Bias

- Defining and characterizing the set of Interesting/Learnable problems
- To what extent do current biases cover the set of interesting problems
- Automatic feature selection
- Automatic selection of Bias (before and/or during learning), including all learning parameters
- Dynamic Inductive Biases (in time and space)
- Combinations of Biases – Ensembles, Oracle Learning


Dynamic Inductive Bias in Time

- Can be discovered as you learn
- May want to learn general rules first followed by true exceptions
- Can be based on ease of learning the problem
- Example: SoftProp – From Lazy Learning to Backprop

Dynamic Inductive Bias in Space



ML Holy Grail: We want all aspects of the learning mechanism automated, including the Inductive Bias



BYU Neural Network and Machine Learning Laboratory

Work on Automatic Discover of Inductive Bias

- Proposing New Learning Algorithms (Inductive Biases)
- Theoretical issues
 - Defining the set of Interesting/Learnable problems
 - Analytical/empirical studies of differences between biases
- Ensembles – Wagging, Mimicking, Oracle Learning, etc.
- Meta-Learning – A priori decision regarding which learning model to use
 - Features of the data set/application
 - Learning from model experience
- Automatic selection of Parameters
 - Constructive Algorithms – ASOCS, DMPx, etc.
 - Learning Parameters – Windowed momentum, Automatic improved distance functions (IVDM)
- Automatic Bias in time – SoftProp
- Automatic Bias in space – Overfitting, sensitivity to complex portions of the space: DMP, higher order features